

Efficient Processing of k Nearest Neighbor Joins using MapReduce

I. SUMMARY

The paper illustrates an efficient processing of kNN Joins using MapReduce.

In the preliminaries, the authors show the basic concepts of kNN Join, MapReduce Framework, and particularly the Voronoi Diagram-based Partitioning, which is the mathematic fundamental for partitioning the input set R and S.

A native and straightforward idea of performing kNN join in MapReduce is similar to the hash join algorithm. R is split into disjoint subsets, each subset R_i is distributed to a reducer, and without any pruning rule, the entire set S has to be sent to each reducer to be joined with R_i .

In order to reduce the shuffling cost, a better strategy is that R is partitioned into N disjoint subsets and for each subset R_i , find a subset of S_i that $R_i \bowtie S = R_i \bowtie S_i$ and $R \bowtie S = \bigcup_{1 \leq i \leq N} R_i \bowtie S_j$. So S_i is sent to the reducer that R_i belongs to and the KNN join is performed between R_i and S_i only.

For the purpose of reducing the size of S_i , we need to derive a distance bound based on the partitioning of R. Using basic geometry knowledge, the upper bound distance from $s \in P_j^S$ to $r \in P_i^R$, denoted as $ub(s, P_i^R)$, can be calculated easily. Then a bound (denoted as θ_i) of the KNN distance for all objects in P_i^R , can be touched by continually adding $ub(s, P_i^R)$, which $s \in KNN(p_j, P_j^S)$, to a priority queue for each P_j^S , until the size grows to k and there is no more $ub(s, P_i^R)$ smaller than the top of the queue.

The corresponding lower bound can be derived in a similar way, and after that, we get the necessary condition s is assigned to S_i :

$$|s, p_j| \geq LB(p_j^S, P_i^R) \quad (1)$$

So, given the input set R and S, the partition of R_i and S_i should be calculated first in the map procedure, and then each R_i, S_i pair will be sent to the same reducer to execute KNN join respectively. The details are shown below, and it takes three steps to complete the kNN Join.

- First, the master node invokes a preprocessing step, which takes the original R and S as input and finds out a set of pivot objects based on the input dataset R. The pivots can be got by Random Selection, Farthest Selection or k-means selection, and the pivots is used to create a Voronoi diagram, which can help partition objects in R effectively while preserving their proximity.
- Second, the first MapReduce job consists of a single Map phase, which takes the selected pivots and datasets R and S as the input. The output of the mapping phase is a partitioning on R, based on the Voronoi diagram of

the pivots. Meanwhile, the mappers also collect some statistics about each partition R_i .

The partitioning on R can be represented as a table, which including each object o along with its partition id, original dataset name (R or S), distance to the closest pivot.

The statistics are kept in two summary table T_R and T_S . T_R maintains the following information for every partition of R: the partition id, the number of objects in the partition, the minimum distance $L(P_i^R)$ and maximum distance $U(P_i^R)$ from an object in partition P_i^R to the pivot. Moreover, T_S also maintains the distances between objects in $KNN(p_i, P_i^S)$ and p_i .

- Third, taking T_R, T_S and statistics from steps 2 as input, S_i will be built with the constraint of equation 1. The R_i and S_i will be sent to the same reducer to execute $R_i \bowtie S_i$. In the reducer, each value pair will be parsed to derive the partition P_i^R and subset S_i that consists of $P_{j1}^S, \dots, P_{jm}^S$. then, a similar pruning process as the second step does, can be applied to S_i to get $LB(P_j^S, P_i^R)$ for every KNN distance, θ_i , for all objects of P_i^R . Hence, we can issue a range search with query and threshold θ_i over dataset S_i . After checking all partitions of S_i with pruning and updating $KNN(r, S)$, the reducer outputs $KNN(r, S)$.

The way S_i derived will bring replications of S. To minimize the number of replicas of objects in S, a intuitive way is to increase the number of pivots. However, this requires a large number of reducers, which may not be practical. A natural idea is to divide partitions of R into disjoint groups. There are two strategies for grouping.

- Geometric Grouping first select N p_i as the basic element of each group G_i , which is faraway from each other, and assign the rest p_j to the nearest group G_i .
- Greedy Grouping tries to minimize the increasing of $RP(S, G_i)$ when assigning a new partition P_j^R to G_i . But for the sake of reducing the computation cost, once $\exists s \in P - l^S$ satisfying the necessity of being added to S_i , we add all objects of partition P_i^S to $RP(S, G_i)$.

The draw back of the algorithm is that, improper pivot selection will cause large difference in partition size, which degrades performance due to unbalanced workload. And this paper doesn't present a pivot selection algorithm that will guarantee the uniformity of partitioning.